

SHREYA ASOBA

AI Engineer | Applied AI | LLM and GenAI Systems

Vancouver, WA (Portland Metro) | Open to relocation | [857-832-1210](tel:857-832-1210) | shreya.asoba@gmail.com
linkedin.com/in/shreya-asoba | github.com/shreyaasoba | [Google Scholar](https://scholar.google.com/citations?user=shreyaasoba)

SUMMARY

AI Engineer with 2.5+ years building, evaluating, and shipping production LLM systems, within 6 years of experience in software, data, and IoT systems. Builds end to end in Go, Python, and TypeScript: model evaluation, AI agents, RAG, document AI pipelines, and full-stack features that ship them to users. Works directly with end users and domain experts to turn real-world failures into fixes.

EXPERIENCE

AI Engineer, Clinical Data | Nephrolytics AI (Healthcare AI) | Remote **Nov 2025 – Present**

- Benchmarked four LLMs for document classification and summarization to replace GPT-5.2, giving each pipeline stage its best performer; cut classification cost 20% and improved accuracy 8%.
- Architected a Go service mapping free-text lab names from EMR APIs and scans to 59K+ LOINC codes at 98% accuracy over 4M+ observations; a prior-mapping cache, 48 skip patterns, and concurrent goroutines keep most records off the LLM path.
- Shipped implausible-value detection across Go, Django, and React: audited every unit and reference range to set per-code bounds, routed impossible OCR values to clinician review, and regression-tested all five downstream AI features.
- Reworked extraction, summarization, and chain-of-thought endnote prompts through recurring sessions with practicing physicians, eliminating failures such as the model silently rewriting raw units (mg% to mg/dL).
- Hardened the document AI pipeline for real-world inputs: PDF, TIFF, and image ingest, image-render failures, a context-window overflow that failed entire batches, and dates in multi-column layouts.
- Diagnosed the accuracy ceiling of AWS Textract and benchmarked LlamaParse via a production Go client behind an interface-compatible extractor; it beat Textract on its failure cases at lower cost (accepted in principle, deferred on budget).

AI & Data Engineer | SQOR.ai (SaaS AI) | New York, NY **Oct 2024 – Mar 2026**

- Built LLM agents with LangGraph, Vertex AI Gemini, and BigQuery that let analysts query multiple data sources in natural language, replacing hand-written SQL for routine analysis.
- Implemented semantic chunking for RAG over financial reports (sentence-transformer embeddings, cosine-similarity clustering, token-aware sizing), improving chunk quality 40% over recursive splitting.
- Built ETL pipelines across 10+ external APIs, expanding metric coverage 20%, and documented 300+ metrics across Snowflake, BigQuery, and SAP, cutting integration research time 60%.

AI Engineer Intern | Reality AI (EdTech AI) | New York, NY **Feb 2024 – Sep 2024**

- Led a team of 8 developers to deliver a notes-generation platform for documents and audio (pdfplumber, Whisper, OpenAI API), designing role-specific system prompts to tailor output for different user groups.
- Built a Vertex AI Gemini Q&A assistant with role-conditioned prompts for students, teachers, and admins (+15% relevance).

AI/ML Instructor | iD Tech, MIT Campus | Boston, MA **Jun 2024 – Aug 2024**

- Taught a 40-hour AI/ML curriculum (CNNs, GANs, LSTMs) to 100+ students, earning 95% satisfaction in post-course surveys.

IT Consultant | Boston University | Boston, MA **Sep 2022 – Dec 2023**

- Built ticket-volume dashboards that cut manual reporting effort 30% while resolving 300+ ServiceNow tickets per month.

Programmer Analyst | Cognizant | Pune, India **Mar 2021 – Aug 2022**

- Built NLP sentiment models, SQL pipelines, and Tableau dashboards on 250K+ support interactions (18% faster responses).

SKILLS

Languages and Frameworks: Go, Python, R, SQL, Django, React, pytest

LLM and GenAI: LLM evaluation, LLM-as-judge, eval harness design, prompt engineering, structured outputs, Claude, Codex, AI agents, MCP, RAG, embeddings, semantic chunking, token and cost telemetry

LLM Frameworks and APIs: LangGraph, LangChain, LlamaParse, OpenAI API, Anthropic API, Vertex AI (Gemini), Whisper

Data and Cloud: PostgreSQL, MongoDB, Snowflake, BigQuery, ETL, AWS (S3, SQS, Textract, AppConfig, Bedrock), GCP, Docker, Kubernetes, Grafana, Tableau

Domain: Document AI, healthcare data (LOINC, EMR, PHI scrubbing, HIPAA-aligned pipelines)

EDUCATION

M.S., Computer Information Systems | Boston University | Jan 2024

B.Tech, Electronics and Communications | Medicaps University | May 2020

PUBLICATIONS AND PROFESSIONAL SERVICE

Peer Reviewer, IEEE InterSYS Conference: reviewed 3 manuscripts on AI production systems ([certificate](#))

Author, IEEE: "Traffic Violation Detection using IoT and Image Processing" ([publication](#))